



Know what your agents can do.

Create or buy Agent Behaviour Policies: for one agent in one deployment, what it can actually reach, what you authorised it to do, and the gap between the two.

ONE PUBLISHED POLICY, COUNTED

Nine things it can do that nobody asked for.

Claude Code on the web, with one repository attached.
Every number decrypted from the vault that describes
it, and the vault is published with its read key.

15

capabilities it can reach

6


that you actually asked for

9

it can do that you didn't ask for

7

with nothing real in the way



Four objects, and only one of them is an opinion.

MEASURED

Reach

Everything the agent can actually access in this deployment, including what nobody thought to check.

ELICITED

Mandate

The job, written down: what the business authorised it to do. The one part only you can supply.

DERIVED

Gap

Reach minus mandate. Never written by hand, and recomputed whenever either side moves.

RECORDED

Barriers

What actually stands in the way of each capability: controls, constraints and open questions.

A behaviour policy carries no score, no rating and no verdict. It describes one agent in one deployment, and what to do about it is the reader's decision.

WHY IT IS NOT A MODEL PROBLEM

The same model is harmless in one setup and serious in another.

Connect it to a mailbox and the narrowest Gmail scope that reads one message reads every message. Nothing about the model changed; the deployment did.

So risk is a property of the deployment rather than of the model, and a record that is not fitted to the deployment it describes is describing somebody else's agent.

WHERE THIS GOES

Three rungs, in order. Each needs the one below it.

AVAILABLE NOW

Describe it

Agent Behaviour Policy. What one agent can reach, what you authorised, the gap, and what stands in the way.

TEMPLATE AVAILABLE

Authorise it

Licence to Operate. The organisation is the authority, the behaviour policy is the instrument, the agent is the licensee.

IN DESIGN

Insure it

Insurability Index. The record an underwriter will accept: scored, dated, with the residual risk owned.

The third rung is the business. The first rung is what is on sale, and it is the one that can be delivered today.

WHAT ARRIVES

An encrypted vault you hold the keys to, not a PDF.

What is in it

The policy and the data behind it. The terms your agent reads, as **AGENTS.md** and **SKILL.md**. A licence-to-operate template, unissued. Every version kept.

What you do with it

Hand a read key to your board, your auditor or your broker and they see exactly what you see. Correct the mandate; the gap recomputes.

Markdown for people, JSON for tools, and nothing in your request path. We never ask for a credential, a token, a key, or access to anything of yours.

FREE, AND PAID

16 published free. Four levels when you want yours.

LEVEL	PRICE	WHEN IT ARRIVES	WHO DOES THE WORK
1 · ABP Pack	£10	immediately	automated
2 · ABP Vault	£50	1 to 2 days	automated
3 · ABP Tailored	£500	1 to 3 days from your reply	a person corrects it
4 · ABP Reviewed	£1,500	1 to 5 days from your reply	two sessions, then a signature

Every level is the same document; what changes is how it arrives and who does the correcting. Prices read from `store.sgit.ai` on 22 September 2026: the store owns the cart, this site owns the material. Before any of it, 16 template policies are published with their read keys, and the prompts to write your own are free.

THE FREE RUNG

Twenty minutes, your own assistant, nothing collected.

Four steps, thirteen prompts, pasted into the assistant you have already connected to your mail. At the end you have a written account of what it can reach, what you meant to authorise, and the gap.

It is the top of the sales motion rather than a giveaway: correcting our draft is how you tell us what you actually intended, and intent is the one input nobody can derive.

Try it

riskmandate.ai/try-it.html

WHO DOES THE WORK

The sign-off carries a name and a date.

At the top level

Two half-hour sessions with your team, the policy built from the interview rather than from a form, and a named security professional who corrects it and signs it.

How the list is built

Every line of a reviewer's page is read off a page they publish themselves, with the date. Nobody is scored, ranked or rated, and we compose nobody's biography.

One name on the list today, and the page is built for more. A second entry is published as a labelled placeholder so the next professional can read exactly what would be said about them before agreeing to it.

WHY THIS YEAR

Two things changed, and neither of them was us.

1 JANUARY 2026

Cover started being withdrawn

A generative-AI exclusion took effect in the standard liability forms much of one large insurance market runs on. Several carriers have filed their own, one of them absolute.

9 MARCH 2026

Responsibility was assigned

A national consumer regulator said a business is responsible if an agent it uses does something illegal, and should be clear what the agent may do and what data it can access.

Both are sourced and dated on riskmandate.ai/insurance.html, with the announcement dates. Neither is our observation, and we are not an insurer and place no cover.

HOW IT IS CHECKABLE

Open source, and the keys are printed.

- **16 behaviour policies published in full**, each with its read key on the page, because a policy nobody can check is a policy asking to be trusted.
- **Every page of this site is in a public repository**, code Apache-2.0, pages and policies CC BY 4.0.
- **The method is written down with its mistakes**: what was measured, what was documented, and what is still an open question, on the page rather than in a footnote.
- **Model-drafted and marked as such**. Not a compliance assessment, and no standards body has assessed anything here.

WHAT WE DO NOT CLAIM

Said before anybody has to ask.

- **We are not an insurer** and we place no cover. The insurance rung is designed, not built.
- **Levels 3 and 4 have never been sold.** They are specified and the payment link is not issued; the store's ledger says so and so does our pricing page.
- **We never test somebody else's system.** Vendor behaviour is read from vendor pages, quoted and dated; where those pages disagree we publish the contradiction unresolved.
- **A behaviour policy reduces accidents. It does not stop an attacker.** A rule somebody wrote down shapes what an agent tends to do, not what it can do.

Start with the agent you cannot describe.

Not the estate. One agent, already running, written down as a record you keep, correct, and can hand to whoever asks.

READ ONE FREE

riskmandate.ai

WRITE YOUR OWN

riskmandate.ai/try-it.html

BUY ONE

store.sgit.ai