



Make every AI agent insurable.

Know what your agents can do. Put a named expert behind the judgement. Hand the evidence to the insurer.

Agent evidence

Human sign-off

Insurability

Nimay Parekh, Co-founder & CEO · Dinis Cruz, Co-founder · riskmandate.ai

The problem

AI agents are becoming economic actors. The liability stack is not ready.

Cover?

Insurers

AI-specific exclusions and underwriting uncertainty are creating a new evidence problem: what exactly did the agent have permission to do?

Accountable

Regulators

Businesses remain accountable for automated actions. Governance must move from model-level policy to deployment-level evidence.

No record

Enterprises

Questionnaires and model cards do not show actual reach, mandate, barriers or residual risk in a live deployment.

The missing primitive is a portable, signed record of what an agent can do — and who stands behind the judgement.

Why now

Three clocks started at once.

\$4.7B

Insurance

Projected annual global AI insurance premiums by 2032. Insurers need better ways to assess and price AI risk.

\$15.8B

Governance

Forrester forecasts off-the-shelf AI governance software spend reaching \$15.8B by 2030.

75%


Regulation

Gartner expects fragmented AI regulation to extend to 75% of world economies by 2030.

AI adoption is outrunning the mechanisms that make autonomous systems governable, auditable and insurable.

The missing piece

The people who can say “yes” are not bookable — and their judgement is not portable.



Turn expert judgement from a conversation into infrastructure.

The insight

Combine continuous evidence with expert judgement — into one insurer-readable record.

Category	Built on	Leaves behind
GRC	Controls	An audit report
Expert networks	Advice	The memory of a call
RiskMandate	Agent evidence + sign-off	A signed, versioned record

RiskMandate is not another dashboard. It is the system of record for an agent’s authority to operate.

The solution

Describe it. Correct it with a named person. Hand it to the insurer.

01

Agent Behaviour Policy

Reach, mandate, gap and barriers for one agent in one deployment.

02

Reviewer bench

Security leaders, former underwriters, examiners, counsel and platform engineers on call.


03

Vault + Licence

Every signed version preserved, shareable by read key, with a time-bound Licence to Operate.

How it works

From model-drafted evidence to human-owned accountability.



The correction is the product. The signature makes the judgement portable.

What the insurer reads

One agent. One deployment. Four questions.

15

Reach

What can it actually access?

6

Mandate

What was it authorised to do?

9

Gap

What is reachable but not authorised?

7

Barriers

What actually prevents each action?

A static model card describes a model. RiskMandate describes the deployed agent an insurer is actually being asked to cover.

The reviewer network

Book the people who can make the judgement credible.

SP Security professional Reach and barriers	UW Former underwriter Reads as the insurer	EX Examiner Regulatory lens
PC AI and privacy counsel Lawful mandate	PE Platform engineer Closes gaps	SB Specialty broker Takes the record to market

Network effect: every reviewer brings enterprise demand; every insurer that accepts a record creates the template for the next deployment.

Product

A wedge today. A new underwriting primitive tomorrow.

1

Available now

Agent Behaviour Policy

Deployment-specific reach, mandate, gap and barriers.

2

Template ready

Licence to Operate

The business is the authority, the policy is the instrument, the agent is the licensee.

3

In design

Insurability Index

A dated, scored record with residual risk explicitly owned.

Long-term: RiskMandate becomes the trust layer between enterprise agents, brokers, carriers and regulators.

Market

We sit at the intersection of three fast-forming budgets.

\$4.7B

AI insurance

Annual global premiums projected by 2032. This is the pool our evidence can help make writable — not our revenue.

\$15.8B

AI governance software

Forrester forecast for off-the-shelf AI governance software spend by 2030.

\$2.5B

Expert networks

Existing expert-network spend proves buyers already pay to access scarce judgement; RiskMandate makes that judgement operational.

Beachhead SAM: enterprises deploying consequential agents in finance, healthcare, commerce, security and operations — where a human sign-off materially changes risk.

Three revenue lines that compound with agent count.

£10–£1.5K

Behaviour policies

Free template → paid deployment policy
→ enterprise licence.

20%*

Sessions + bench

Marketplace take rate on expert sessions
plus recurring standing-bench
subscriptions.

£500+*

Insurability Index

Per-agent or portfolio licence sold to
enterprises, brokers and insurers.

*Proposed launch economics for investor discussion; validate with reviewer supply and insurer design partners.

Competition

Everyone owns one piece of the loop.
We connect the loop.

Category	Deployment evidence	Named sign-off	Insurer-readable
GRC automation	Controls, not agents	No	Limited
AI governance	Mostly model or system level	No	Limited
AI insurers / MGAs	For their own book	No	Own underwriting
Expert networks	No	Yes	No
RiskMandate	Yes	Yes	Yes

Neutral by design: RiskMandate does not underwrite or place cover. It makes judgement portable across the market.

Land with open evidence. Expand through the people who already hold trust.

1

Land

16 open behaviour-policy templates let teams see their own agent shape before buying.

2

Expand

The policy exposes the gap. The gap creates the reviewer session. The reviewer becomes the standing bench.

3

Network

Brokers, reviewers and agent platforms become distribution because each needs a reusable evidence standard.

First design-partner motion: specialty brokers + AI-forward carriers + agent platforms + CISOs governing production agents.

Traction

Live product, open standard, real artefacts — before institutional capital.

16

Templates

Behaviour-policy templates published under CC BY 4.0.

Live

Working vault

Customer-held keys and a reading app already exist.

Shipped

Review flow

Tailored and Reviewed levels delivered by a named person.

Next proof points to unlock: 10+ committed reviewers · 3 broker or carrier design partners · first recurring enterprise bench.

These are forward milestones, not claimed current traction.

Team

Cyber risk scoring meets security engineering and GenAI.

Nimay Parekh

Co-founder & CEO

Built cyber-insurance and supply-chain risk products at SecurityScorecard.

Senior partnerships and product leadership across enterprise software; INSEAD MBA; White House innovation fellowship experience.

Dinis Cruz

Co-founder


Long-time application-security and OWASP leader.

Former CISO, Holland & Barrett. Chief Scientist, Glasswall. Founder, The Cyber Boardroom. Deep focus on GenAI, security and operational risk.

Founder-market fit: enterprise risk + insurance distribution + security engineering + AI-native governance.

The vision

Every consequential agent should have a licence to operate.



From “trust me” to a verifiable chain of authority: what the agent can do, what it may do, what stops it, and who signed.

The trust layer for the agentic economy.

Evidence that enterprises can own. Judgement that experts can sign. Records that insurers can read.

riskmandate.ai

Nimay Parekh • Dinis Cruz