

RISKMANDATE.AI

Who can pull the plug?

Every board asks it. Almost no enterprise can answer it for its agents.

We make the answer computable — and make someone own it.

Never in line

Open core

Consumption-priced

Enterprises cannot say who stops their agents

“If this agent starts doing something no one intended — who stops it, how fast, and what does stopping it cost?”

~0

agents with a stated plug profile

the answer is a pause, not a name

12–18h

typical detection floor

exposure before anyone sees it

100s

agent actions per human checkpoint

the grace period is gone

The plug always exists. What no one can state is its profile — who, blast radius, speed, side effects, and how much is recoverable.

We compute the plug profile — and get it accepted

Computed, not asserted

Every field derives from the graph and carries evidence — not a workshop opinion or a colour on a heat map.

Recoverability decides

The one dimension money cannot buy back. It separates a catastrophic-but-recoverable risk from a smaller irreversible one.

An acceptance, not a record

A named owner, conditions, an expiry. The profile is the artifact the board actually wants.

PLUG PROFILE
AP-Agent-07 · invoice approval

WHO	S. Okafor (Finance Ops)	authority + capability verified
BLAST RADIUS	AP pipeline + 3 downstream	\$1.4M exposure
SPEED	4 min to pull	damage doubles every 9 min
SIDE EFFECTS	812 in-flight transactions	recoverable
RECOVERABILITY	92% recoverable	8% irreversible: sent payments

ACCEPTED S. Okafor · expires 30 Sep 2026 · 6 evidence artifacts

Five dimensions. One owner. An expiry date.

THE CORE MECHANIC

You don't get to not accept the risk

Ask someone to accept a risk and they will do neither. Treat acceptance as a choice and you hand them a third door: do nothing.

~~Do you accept this risk? Yes / No / Later~~

The risk is already on your books. Acceptance is not a decision to take it on — it is an acknowledgement that you already have.

DECISION 01 · DURATION

For how long?

Every acceptance expires. When it does, the same decision lands back on the same desk — with the risk still there.

30d

90d

180d

DECISION 02 · NEXT ACTION

And what happens next?

Either you buy more detail, or you buy the fix. Both cost money, and both are choices someone's name goes against.

Investigate

Remediate

No third door. Silence is not an available answer.

The Act mandates the judgement — not the word

EU AI Act Article 9(5) requires that residual risk be judged acceptable. The obligation is imposed; the standard is not supplied.

THE GAP THE ACT LEAVES OPEN

You must decide your line.
You must show you decided.

Not having defined it is itself a first-class risk — and it is the first question we ask.

DECLARED vs REVEALED APPETITE

The stated line, versus where remediation actually stopped.

The gap between the two is the finding. It is answerable without tooling — and almost never answered.

INTERVAL IS COMPUTED FROM DISTANCE TO THE LINE

Far above the line

~1 month

high gap · short interval

Mid above the line

~3 months

medium gap

Near or at the line

~6 months


small gap · long interval

Illustrative intervals. Distance sets the clock — the further above the line, the sooner it returns to the desk.

THE DISTINCTION


Accepted is not acceptable

One is an act. The other is a line. They are orthogonal — which gives four real states, each needing a different action.



Missing acceptance is an ownership problem. Missing acceptable line is a governance problem. Different failures, different fixes.

Read-only by design — twins instead of integrations



NEVER IN LINE — no runtime decisions, no enforcement point, no request path

No integration treadmill

Twins model the primitives that matter — assets and flows — so we reach every system without connecting to any.

Never in the request path


No runtime decisions, no enforcement. Security review and procurement collapse from months to weeks.

Telemetry tests the map

Signals are consumed to detect mapping error, never to police traffic. Capability is reverse-engineered from behaviour.

Authorization is whatever the agent can already do

Not a decision at request time. The moment capability was conferred — measured as what is actually possible.



= THE UNION OF EVERYTHING THE AGENT CAN ACTUALLY DO

If the agent can do it, it was authorised.

The agent that deleted the database was, by our definition, already authorised to.

Our value is shrinking it.

Make the exposure visible, bound it, and cut exploitation and collateral damage.

This is why we never need to be in line — authorization already happened.

A flat register cannot survive the interrogation

GRC / ORM

MetricStream · Archer · ServiceNow

- Records risk
- Flat register
- Scores, asserted
- Overhead

IDENTITY / CSPM

CyberArk · Okta · Wiz

- Grants access
- Finds misconfig
- No accountability
- No acceptance

RISKMANDATE

the decision instrument

- **Computes the plug profile**
- **Fractal, every altitude**
- **Evidence, not scores**
- **Accepted, owned, funded**

Acceptance transfers accountability to a named human. That triggers demands for evidence, accuracy and correct attribution — which a scored, flat register structurally cannot meet.

Two KPIs the board can read

KPI 1

Risks accepted

Decisions actually made — a named owner, conditions, an expiry. Volume weighted by exposure, so trivia cannot game it.

KPI 2

Business risk reduction

Exposure removed, measured against a stated baseline. The number the board is asking for.

OPERATING METRICS THAT MOVE FIRST · ILLUSTRATIVE TARGETS

METRIC	BEFORE	WITH RISKMANDATE
Agents with a stated plug profile	12%	94%
Mean time to risk acceptance	41 days	6 days
Risks with a funded owner	23%	87%
Irreversible exposure quantified	not known	100% of critical agents

Governance stops being an assertion and becomes a measured number.

Consumption, priced like a hyperscaler

De-risks adoption

No seat licence, no upfront commitment. The platform enters without a procurement fight.

Hook is delivered value

Lock-in comes from what the graph gives back — never from the contract.

Customer can cap it

Spend is bounded at the buyer's discretion, which answers the budget objection directly.

Aligns both sides

We and the customer are rewarded for using the platform more — exactly the hyperscaler dynamic.

METERED DIMENSIONS

Compute

Tokens

Bandwidth

Storage

Access

Uptime

Not a seat licence.

Revenue scales with twins modelled, agents profiled and briefings produced — usage, not headcount.

Caps trade a revenue ceiling for adoption speed — a deliberate choice.

Self-serve the map. Sell the decision.

SELF-SERVE

The map

- Open-source core + twins
- Credit card, no procurement
- Minutes to first exposure
- Metered, cappable

Near-zero CAC



HUMAN-LED

The decision

- Risk acceptance itself
- A named owner takes accountability
- Evidence · accuracy · attribution
- Expands by altitude

High ACV, sales-assisted

The handoff: open source builds the map for free, and the map creates the acceptance conversation. The first exposure a customer cannot get accepted is the qualified lead — and “Who can pull the plug?” is the content engine that produces it.